

Single Logistic Regression (Binary classification)

import required package

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

read data from source and describing

```
In [2]: df = pd.read_csv('social_network_ads.csv')
df.head()
```

```
Out[2]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [3]: print(df.columns)
```

Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'], dtype='object')

```
In [4]: print(df.describe())
```

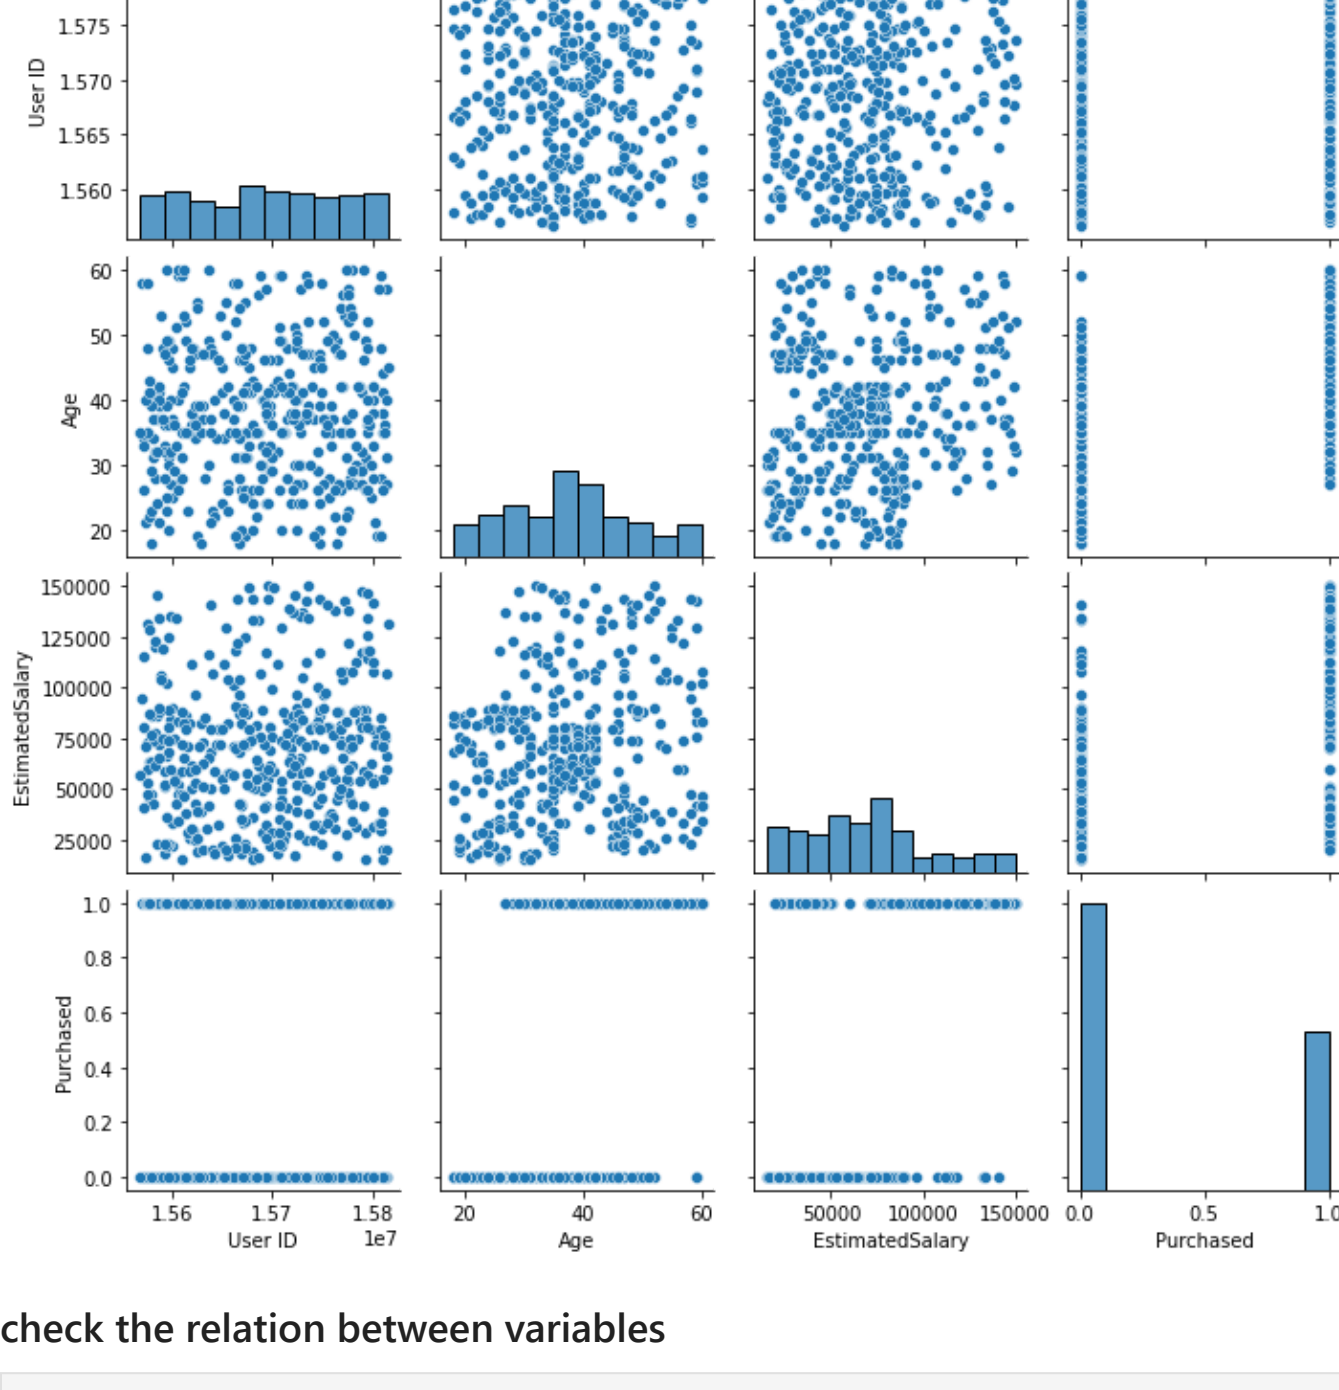
```
count      User ID      Age  EstimatedSalary  Purchased
mean  4.000000e+02  37.655000      69742.500000  0.357500
std    7.165832e+04  10.482877      34096.960282  0.479864
min    1.556669e+07  18.000000      15000.000000  0.000000
25%    1.562676e+07  29.750000      43000.000000  0.000000
50%    1.569434e+07  37.000000      70000.000000  0.000000
75%    1.575036e+07  46.000000      88000.000000  1.000000
max    1.581524e+07  60.000000     150000.000000  1.000000
```

```
In [5]: ## chec for data types
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column             Non-Null Count  Dtype
---  --
0   User ID             400 non-null   int64
1   Gender              400 non-null   object
2   Age                 400 non-null   int64
3   EstimatedSalary     400 non-null   int64
4   Purchased           400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
None
```

```
In [6]: sns.pairplot(df)
```

Out[6]: <seaborn.axisgrid.PairGrid at 0x1a09d44d160>



check the relation between variables

```
In [7]: corr = df.corr()
corr.style.background_gradient(cmap = 'Greens_r')
```

```
Out[7]:
```

	User ID	Age	EstimatedSalary	Purchased
User ID	1.000000	-0.000721	0.071097	0.007120
Age	-0.000721	1.000000	0.155238	0.622454
EstimatedSalary	0.071097	0.155238	1.000000	0.362083
Purchased	0.007120	0.622454	0.362083	1.000000

```
In [8]: ## normalization of data
#convert gender to unique values
vals = df['Gender'].unique()
print(vals)

df.replace(vals, [0, 1], inplace= True)
# print(df['Gender'])
print(df.head())
```

```
['Male' 'Female']
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510      0   19         19000          0
1  15810944      0   35         20000          0
2  15668575      1   26         43000          0
3  15603246      1   27         57000          0
4  15804002      0   19         76000          0
```

```
In [9]: ### check the relation between variables
```

```
corr = df.corr()
corr.style.background_gradient(cmap = 'Greens')
```

```
Out[9]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
User ID	1.000000	0.025249	-0.000721	0.071097	0.007120
Gender	0.025249	1.000000	0.073741	0.060435	0.042469
Age	-0.000721	0.073741	1.000000	0.155238	0.622454
EstimatedSalary	0.071097	0.060435	0.155238	1.000000	0.362083
Purchased	0.007120	0.042469	0.622454	0.362083	1.000000

using ordinal encoder

```
In [10]: # from sklearn.preprocessing import OrdinalEncoder
encoder = OrdinalEncoder()
encoder.fit(df[["Gender"]])
df[["Gender"]] = encoder.transform(df[["Gender"]])
# print(df)
```

select input and op variable

```
In [11]: x = df.drop(['User ID', 'Gender', 'Purchased'], axis=1)
y = df['Purchased']
```

split the data

```
In [12]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state= 123456, train
```

creating a model

```
In [13]: from sklearn.linear_model import LogisticRegressionCV
model = LogisticRegressionCV()
```

```
In [14]: ## fit the data
model.fit(x_train, y_train)
```

Out[14]: LogisticRegressionCV()

parameters to fine tune model

```
In [15]: print(model.get_params())
```

```
{'Cs': 10, 'class_weight': None, 'cv': None, 'dual': False, 'fit_intercept': True, 'in
tercept_scaling': 1.0, 'l1_ratios': None, 'max_iter': 100, 'multi_class': 'auto', 'n_j
obs': None, 'penalty': 'l2', 'random_state': None, 'refit': True, 'scoring': None, 'so
lver': 'lbfgs', 'tol': 0.0001, 'verbose': 0}
```

predict the values for unseen data

```
In [16]: y_prediction = model.predict(x_test)
#print(y_prediction)
```

evaluation of classification model

```
In [17]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_prediction)
print(cm)
```

```
[[55  1]
 [10 14]]
```

```
In [18]: correct = cm[0,0] + cm[1,1]
wrong = cm[1,0] + cm[0,1]
total = correct + wrong
accuracy = correct/total
print(accuracy)
```

0.8625

```
In [19]: from sklearn.metrics import accuracy_score
print(accuracy_score(y_test, y_prediction))
```

0.8625

precision value

```
In [20]: from sklearn.metrics import precision_score
print(precision_score(y_test, y_prediction))
```

0.9333333333333333

recal value

```
In [21]: from sklearn.metrics import recall_score
print(recall_score(y_test, y_prediction))
```

0.5833333333333334

F1 score

```
In [22]: from sklearn.metrics import f1_score
print(f1_score(y_test, y_prediction))
```

0.7179487179487181

classification report

```
In [23]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_prediction))
```

```
precision    recall  f1-score   support
```

```
0           0.85     0.98     0.91         56
```

```
1           0.93     0.58     0.72         24
```

```
accuracy          0.86         80
```

```
macro avg         0.89     0.78     0.81         80
```

```
weighted avg      0.87     0.86     0.85         80
```

```
In [24]: from sklearn.metrics import roc_curve, roc_auc_score, plot_roc_curve
```

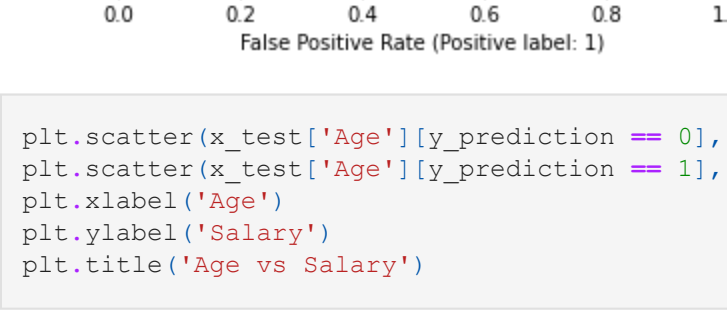
```
print(roc_auc_score(y_test, y_prediction))
```

0.7827380952380953

```
In [25]: fpr, tpr, threshold = roc_curve(y_test, y_prediction)
print(fpr)
print(tpr)
print(threshold)
plt.plot(fpr, tpr)
```


```
[0.         0.01785714 1.         ]
[0.         0.58333333 1.         ]
[2 1 0]
```

Out[25]: <matplotlib.lines.Line2D at 0x1a09f80beb0>



```
In [26]: plot = plot_roc_curve(model, x_test, y_test)
plt.plot([0,1], [0,1], linestyle = '--')
```

Out[26]: <matplotlib.lines.Line2D at 0x1a09f871e80>



In []: