

Multiple Linear Regression

import required package

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

read data from source and describing

```
In [2]: df = pd.read_csv('50_Startups.csv')
df.head()
```

```
Out[2]:
```

	RnD	Administration	Marketing	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
In [3]: print(df.columns)

Index(['RnD', 'Administration', 'Marketing', 'State', 'Profit'], dtype='object')
```

```
In [4]: print(df.describe())
```

	RnD	Administration	Marketing	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	112012.639200
std	45902.256482	28017.802755	122290.310726	40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	90138.902500
50%	73051.080000	122699.795000	212716.240000	107978.190000
75%	101602.800000	144842.180000	299469.085000	139765.977500
max	165349.200000	182645.560000	471784.100000	192261.830000

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  ------                -
0   RnD                    50 non-null    float64
1   Administration         50 non-null    float64
2   Marketing              50 non-null    float64
3   State                  50 non-null    object
4   Profit                 50 non-null    float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

check the relation between variables

```
In [6]: ## state value is excluded as it is categorical
corr = df.corr()
corr.style.background_gradient(cmap = 'Greens')
```

```
Out[6]:
```

	RnD	Administration	Marketing	Profit
RnD	1.000000	0.241955	0.724248	0.972900
Administration	0.241955	1.000000	-0.032154	0.200717
Marketing	0.724248	-0.032154	1.000000	0.747766
Profit	0.972900	0.200717	0.747766	1.000000

data cleansing and normilisation

```
In [7]: ## since state values are categorical it would be replaced with other values
```

```
In [8]: state_values = df['State'].unique()
print(state_values)

unique_values = [1,2,3]

df.replace(state_values, unique_values, inplace=True)
print(df.head())

['New York' 'California' 'Florida']
```

	RnD	Administration	Marketing	State	Profit
0	165349.20	136897.80	471784.10	1	192261.83
1	162597.70	151377.59	443898.53	2	191792.06
2	153441.51	101145.55	407934.54	3	191050.39
3	144372.41	118671.85	383199.62	1	182901.99
4	142107.34	91391.77	366168.42	3	166187.94

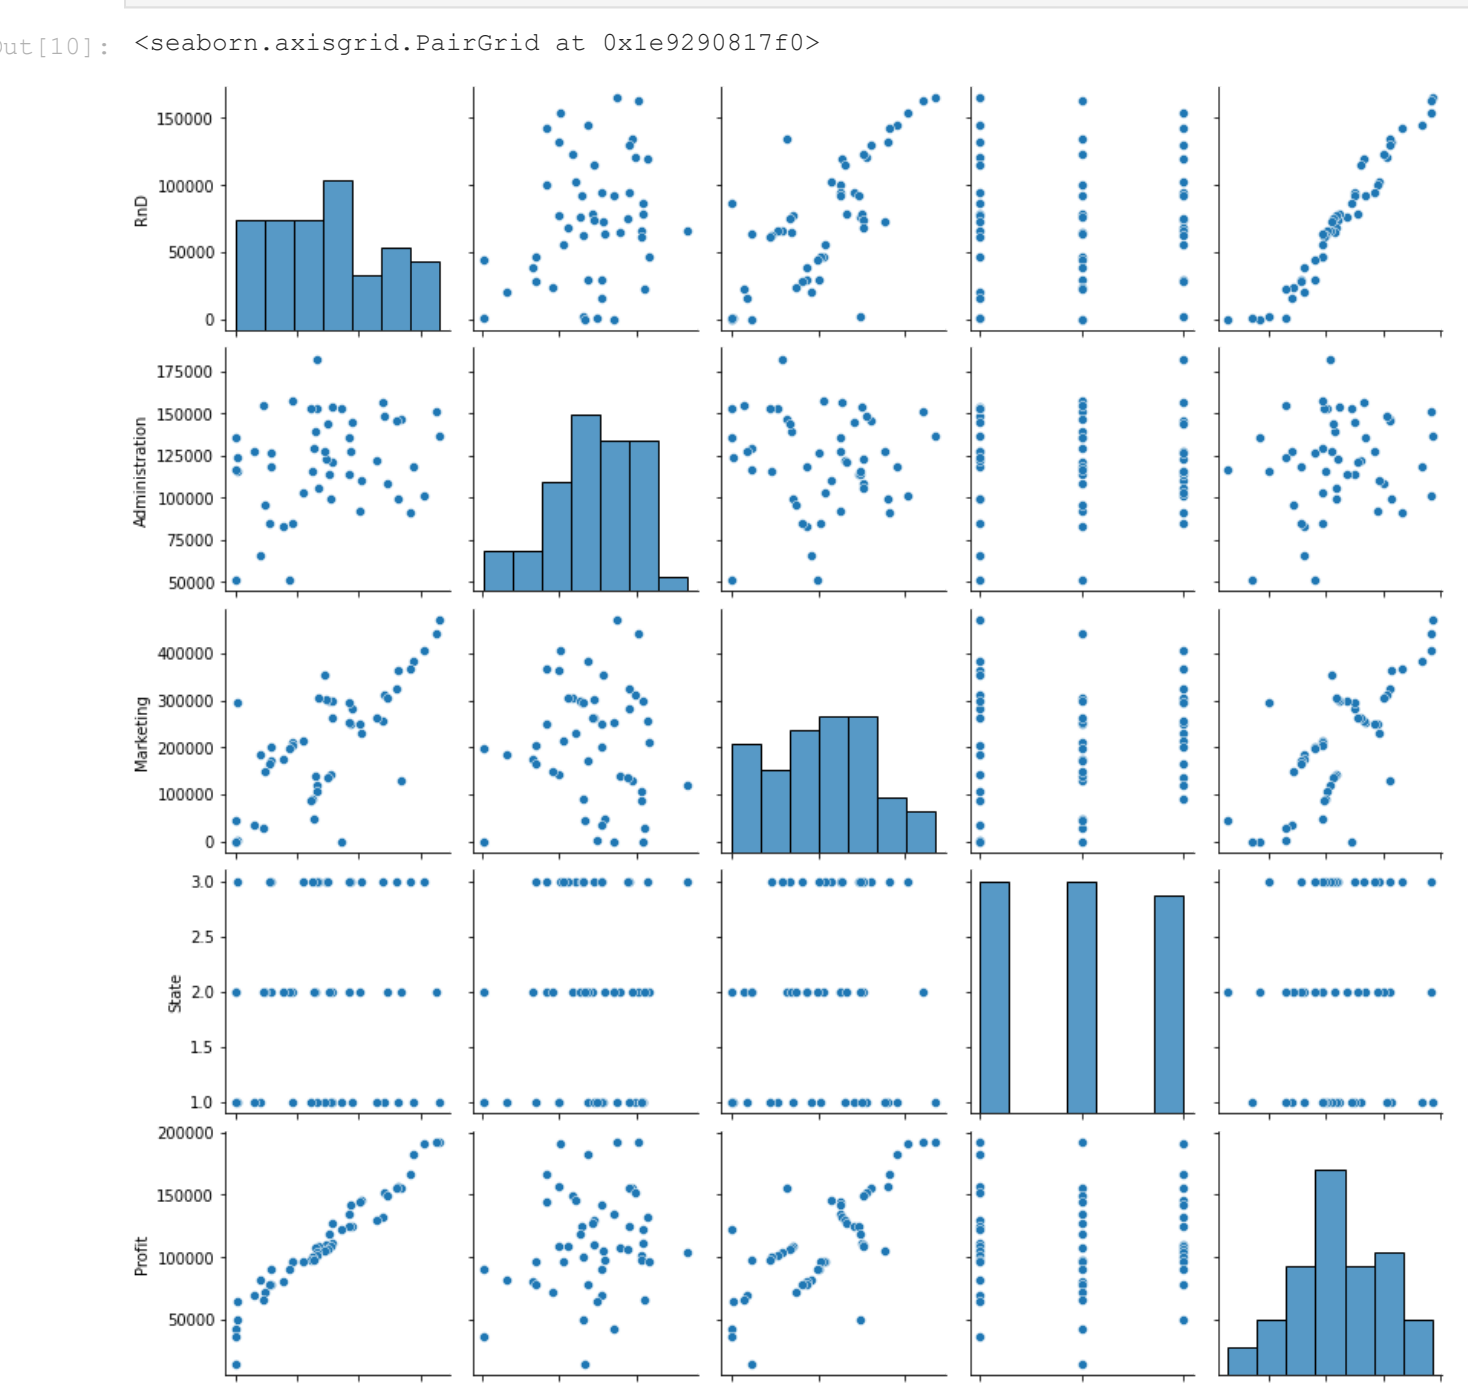
EDA

```
In [9]: corr = df.corr()
corr.style.background_gradient(cmap = 'Greens')
```

```
Out[9]:
```

	RnD	Administration	Marketing	State	Profit
RnD	1.000000	0.241955	0.724248	0.037930	0.972900
Administration	0.241955	1.000000	-0.032154	0.003026	0.200717
Marketing	0.724248	-0.032154	1.000000	0.137777	0.747766
State	0.037930	0.003026	0.137777	1.000000	0.048471
Profit	0.972900	0.200717	0.747766	0.048471	1.000000

```
In [10]: sns.pairplot(df)
```



select input and op variable

```
In [11]: x = df.drop(['State', 'Profit'], axis=1)
y = df['Profit']
```

split the data

```
In [12]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=123456, train_size=0.8)
```

creating a model

```
In [13]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

```
In [14]: ## fit the data
model.fit(x_train, y_train)
```

```
Out[14]: LinearRegression()
```

evaluate the model

```
In [15]: score = model.score(x_train, y_train)
print(score)
```

0.9601722295771402

parameters to fine tune model

```
In [16]: print(model.get_params())
```

{'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'normalize': False, 'positive': False}

predict the values for unseen data

```
In [17]: y_prediction = model.predict(x_test)
print(y_prediction)
```

[67329.47406593 117366.3802916 91360.40144635 154304.50169146
 76117.37130998 98755.52129056 50827.79389197 99500.68282807
 129036.46220559 103456.53045388]

evaluation of loss functions

```
In [18]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
MAE = mean_absolute_error(y_test, y_prediction)
MSE = mean_squared_error(y_test, y_prediction)
RMSE = np.sqrt(MSE)
R2 = r2_score(y_test, y_prediction)

print(MAE)
print(MSE)
print(RMSE)
print(R2)
```

7046.237337953069
149018716.7626929
12207.322260131126
0.8813666867792831

visualization only possible for 2d values

```
In [ ]:
```